

Final Project Submission

Please fill out:

- Student name: Dapash Kumati
- Student pace: full time
- Scheduled project review date/time: 3/10/2025
- Instructor name: Nikita Njoroge
- Blog post URL:

1: Background

The company is expanding its portfolio by entering the aviation industry, focusing on both commercial and private air travel. However, aviation carries inherent safety and financial risks, especially when it comes to aircraft accidents, reliability, and maintenance. To ensure a successful and sustainable entry into this industry, the company requires an evidence-based assessment of aircraft safety and associated risks.

2: Project Overview

This project leverages historical accident data from the National Transportation Safety Board (NTSB), covering civil aviation accidents and selected incidents from 1962 to 2023. By analyzing this dataset, we aim to identify aircraft types with the lowest safety risks. The insights generated will guide business stakeholders in making informed decisions on which aircraft are most suitable for purchase and operation, minimizing risk exposure and maximizing operational reliability.

3: Business Understanding

Goal: Determine which aircraft types are associated with the lowest risk for accidents/incidents.

Objectives:

1. To assess the distribution and severity of aircraft accidents across different aircraft models and manufacturers.
2. To analyze the relationship between accident characteristics (fatalities, damage level, risk category) and the likelihood of aircraft being involved in fatal incidents.

3. To identify trends in aviation accident frequency and severity over time and across geographical locations.

4: Data Understanding

a).Set up the environment and import libraries

```
In [1]: # Basic Libraries
import pandas as pd
import numpy as np
# Visualization Libraries
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
# Display settings
pd.set_option("display.max_columns", None)
```

b).Load Data

```
In [2]: # Load dataset
df = pd.read_csv("data/Aviation_Data.csv", low_memory= False)
df.head()
```

```
Out[2]:
```

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Cou
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	U S
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	U S
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	U S
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	U S
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	U S

c). Initial Data Exploration

```
In [3]: # Shape
df.shape
```

```
Out[3]: (90348, 31)
```

```
In [4]: # Basic information
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 90348 entries, 0 to 90347
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             88889 non-null  object
1   Investigation.Type                    90348 non-null  object
2   Accident.Number                      88889 non-null  object
3   Event.Date                           88889 non-null  object
4   Location                             88837 non-null  object
5   Country                             88663 non-null  object
6   Latitude                             34382 non-null  object
7   Longitude                             34373 non-null  object
8   Airport.Code                         50249 non-null  object
9   Airport.Name                         52790 non-null  object
10  Injury.Severity                      87889 non-null  object
11  Aircraft.damage                      85695 non-null  object
12  Aircraft.Category                    32287 non-null  object
13  Registration.Number                  87572 non-null  object
14  Make                                 88826 non-null  object
15  Model                               88797 non-null  object
16  Amateur.Built                       88787 non-null  object
17  Number.of.Engines                    82805 non-null  float64
18  Engine.Type                          81812 non-null  object
19  FAR.Description                      32023 non-null  object
20  Schedule                             12582 non-null  object
21  Purpose.of.flight                    82697 non-null  object
22  Air.carrier                          16648 non-null  object
23  Total.Fatal.Injuries                 77488 non-null  float64
24  Total.Serious.Injuries               76379 non-null  float64
25  Total.Minor.Injuries                 76956 non-null  float64
26  Total.Uninjured                      82977 non-null  float64
27  Weather.Condition                    84397 non-null  object
28  Broad.phase.of.flight                61724 non-null  object
29  Report.Status                        82508 non-null  object
30  Publication.Date                     73659 non-null  object
dtypes: float64(5), object(26)
memory usage: 21.4+ MB

```

```

In [5]: # Summary statistics for all columns
df.describe(include="all").T

```


Out[5]:

	count	unique	top	freq	mean	std	min	max
Event.Id	88889	87951	20001212X19172	3	NaN	NaN	NaN	1
Investigation.Type	90348	71	Accident	85015	NaN	NaN	NaN	1
Accident.Number	88889	88863	DCA22LA135	2	NaN	NaN	NaN	1
Event.Date	88889	14782	2000-07-08	25	NaN	NaN	NaN	1
Location	88837	27758	ANCHORAGE, AK	434	NaN	NaN	NaN	1
Country	88663	219	United States	82248	NaN	NaN	NaN	1
Latitude	34382	25589	332739N	19	NaN	NaN	NaN	1
Longitude	34373	27154	0112457W	24	NaN	NaN	NaN	1
Airport.Code	50249	10375	NONE	1488	NaN	NaN	NaN	1
Airport.Name	52790	24871	Private	240	NaN	NaN	NaN	1
Injury.Severity	87889	109	Non-Fatal	67357	NaN	NaN	NaN	1
Aircraft.damage	85695	4	Substantial	64148	NaN	NaN	NaN	1
Aircraft.Category	32287	15	Airplane	27617	NaN	NaN	NaN	1
Registration.Number	87572	79105	NONE	344	NaN	NaN	NaN	1
Make	88826	8237	Cessna	22227	NaN	NaN	NaN	1
Model	88797	12318	152	2367	NaN	NaN	NaN	1
Amateur.Built	88787	2	No	80312	NaN	NaN	NaN	1
Number.of.Engines	82805	NaN	NaN	NaN	1.14659	0.44651	0	
Engine.Type	81812	13	Reciprocating	69530	NaN	NaN	NaN	1
FAR.Description	32023	31	091	18221	NaN	NaN	NaN	1
Schedule	12582	3	NSCH	4474	NaN	NaN	NaN	1
Purpose.of.flight	82697	26	Personal	49448	NaN	NaN	NaN	1
Air.carrier	16648	13590	Pilot	258	NaN	NaN	NaN	1
Total.Fatal.Injuries	77488	NaN	NaN	NaN	0.647855	5.48596	0	
Total.Serious.Injuries	76379	NaN	NaN	NaN	0.279881	1.54408	0	
Total.Minor.Injuries	76956	NaN	NaN	NaN	0.357061	2.23563	0	
Total.Uninjured	82977	NaN	NaN	NaN	5.32544	27.9136	0	
Weather.Condition	84397	4	VMC	77303	NaN	NaN	NaN	1
Broad.phase.of.flight	61724	12	Landing	15428	NaN	NaN	NaN	1
Report.Status	82508	17007	Probable Cause	61754	NaN	NaN	NaN	1
Publication.Date	73659	2923	25-09-2020	16317	NaN	NaN	NaN	1

```
In [6]: # Summary statistics for numerical columns
df.describe()
```

```
Out[6]:
```

	Number.ofEngines	Total.Fatal.Injuries	Total.Serious.Injuries	Total.Minor.Injuries
count	82805.000000	77488.000000	76379.000000	76956.000000
mean	1.146585	0.647855	0.279881	0.357061
std	0.446510	5.485960	1.544084	2.235625
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000	0.000000
50%	1.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	0.000000	0.000000
max	8.000000	349.000000	161.000000	380.000000



```
In [7]: df.columns
```

```
Out[7]: Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date',
              'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Code',
              'Airport.Name', 'Injury.Severity', 'Aircraft.damage',
              'Aircraft.Category', 'Registration.Number', 'Make', 'Model',
              'Amateur.Built', 'Number.of.Engines', 'Engine.Type', 'FAR.Description',
              'Schedule', 'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries',
              'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured',
              'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',
              'Publication.Date'],
              dtype='object')
```

5: Data Cleaning

A). Checking missing values

```
In [8]: df.isna().sum().sort_values(ascending=False)
```

```
Out[8]: Schedule      77766
Air.carrier          73700
FAR.Description      58325
Aircraft.Category    58061
Longitude            55975
Latitude             55966
Airport.Code         40099
Airport.Name         37558
Broad.phase.of.flight 28624
Publication.Date     16689
Total.Serious.Injuries 13969
Total.Minor.Injuries 13392
Total.Fatal.Injuries 12860
Engine.Type          8536
Report.Status        7840
Purpose.of.flight    7651
Number.of.Engines    7543
Total.Uninjured      7371
Weather.Condition    5951
Aircraft.damage      4653
Registration.Number  2776
Injury.Severity      2459
Country              1685
Amateur.Built        1561
Model                1551
Make                 1522
Location             1511
Event.Date           1459
Accident.Number      1459
Event.Id             1459
Investigation.Type    0
dtype: int64
```

```
In [9]: # To know the percentage of the missing values
df.isna().mean()*100
```

```
Out[9]: Event.Id          1.614867
Investigation.Type      0.000000
Accident.Number        1.614867
Event.Date             1.614867
Location               1.672422
Country                1.865011
Latitude               61.944924
Longitude              61.954886
Airport.Code           44.382831
Airport.Name           41.570372
Injury.Severity        2.721698
Aircraft.damage        5.150086
Aircraft.Category      64.263736
Registration.Number    3.072564
Make                  1.684597
Model                  1.716695
Amateur.Built          1.727764
Number.of.Engines      8.348829
Engine.Type            9.447913
FAR.Description        64.555939
Schedule               86.073848
Purpose.of.flight      8.468367
Air.carrier            81.573471
Total.Fatal.Injuries   14.233851
Total.Serious.Injuries 15.461327
Total.Minor.Injuries   14.822686
Total.Uninjured        8.158454
Weather.Condition      6.586753
Broad.phase.of.flight  31.681941
Report.Status          8.677558
Publication.Date       18.471909
dtype: float64
```

B). Handling missing values

```
In [10]: # Drop columns with too many missing values (>60% missing)
threshold = 0.6 * len(df)
df = df.dropna(axis=1, thresh=threshold)
```

```
In [11]: # Fill missing categorical values with "Unknown"
for col in df.select_dtypes(include="object").columns:
    df[col] = df[col].fillna("Unknown")
```

```
In [12]: for col in df.select_dtypes(include=["float64", "int64"]).columns:
    df[col] = df[col].fillna(df[col].median())
```

```
In [13]: df.isna().sum()
```

```
Out[13]: Event.Id      0
Investigation.Type    0
Accident.Number       0
Event.Date            0
Location              0
Country               0
Injury.Severity       0
Aircraft.damage       0
Registration.Number    0
Make                  0
Model                 0
Amateur.Built         0
Number.ofEngines      0
Engine.Type           0
Purpose.of.flight     0
Total.Fatal.Injuries  0
Total.Serious.Injuries 0
Total.Minor.Injuries  0
Total.Uninjured       0
Weather.Condition     0
Broad.phase.of.flight 0
Report.Status         0
Publication.Date      0
dtype: int64
```

```
In [14]: df.shape
```

```
Out[14]: (90348, 23)
```

```
In [15]: df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 90348 entries, 0 to 90347
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             90348 non-null  object
1   Investigation.Type                    90348 non-null  object
2   Accident.Number                      90348 non-null  object
3   Event.Date                           90348 non-null  object
4   Location                             90348 non-null  object
5   Country                             90348 non-null  object
6   Injury.Severity                      90348 non-null  object
7   Aircraft.damage                      90348 non-null  object
8   Registration.Number                  90348 non-null  object
9   Make                                90348 non-null  object
10  Model                                90348 non-null  object
11  Amateur.Built                        90348 non-null  object
12  Number.of.Engines                    90348 non-null  float64
13  Engine.Type                          90348 non-null  object
14  Purpose.of.flight                    90348 non-null  object
15  Total.Fatal.Injuries                  90348 non-null  float64
16  Total.Serious.Injuries                90348 non-null  float64
17  Total.Minor.Injuries                  90348 non-null  float64
18  Total.Uninjured                       90348 non-null  float64
19  Weather.Condition                    90348 non-null  object
20  Broad.phase.of.flight                 90348 non-null  object
21  Report.Status                         90348 non-null  object
22  Publication.Date                      90348 non-null  object
dtypes: float64(5), object(18)
memory usage: 15.9+ MB
```

C). Checking and dropping duplicates

```
In [16]: df.duplicated().sum()
```

```
Out[16]: 1390
```

```
In [17]: df = df.drop_duplicates()
print("After dropping full duplicates:", df.shape)
```

After dropping full duplicates: (88958, 23)

```
In [18]: df.duplicated().sum()
```

```
Out[18]: 0
```

D). Outliers

a) Focus on numerical columns

```
In [19]: # Select numeric columns
numeric_cols = df.select_dtypes(include=["float64", "int64"]).columns
print("Numeric columns:", numeric_cols)
```


```
Numeric columns: Index(['Number.of.Engines', 'Total.Fatal.Injuries', 'Total.Serious.Injuries',
                        'Total.Minor.Injuries', 'Total.Uninjured'],
                        dtype='object')
```

b) Visualization to show outliers

Boxplots for Outlier Visualization - are the best way to see outliers (they appear as dots beyond the whiskers).

```
In [20]: import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(12, 6))
sns.boxplot(data=df[numeric_cols], orient="h", palette="Set2")
plt.title("Boxplot of Numeric Columns (Outliers Visualized)", fontsize=14)
plt.show()
```

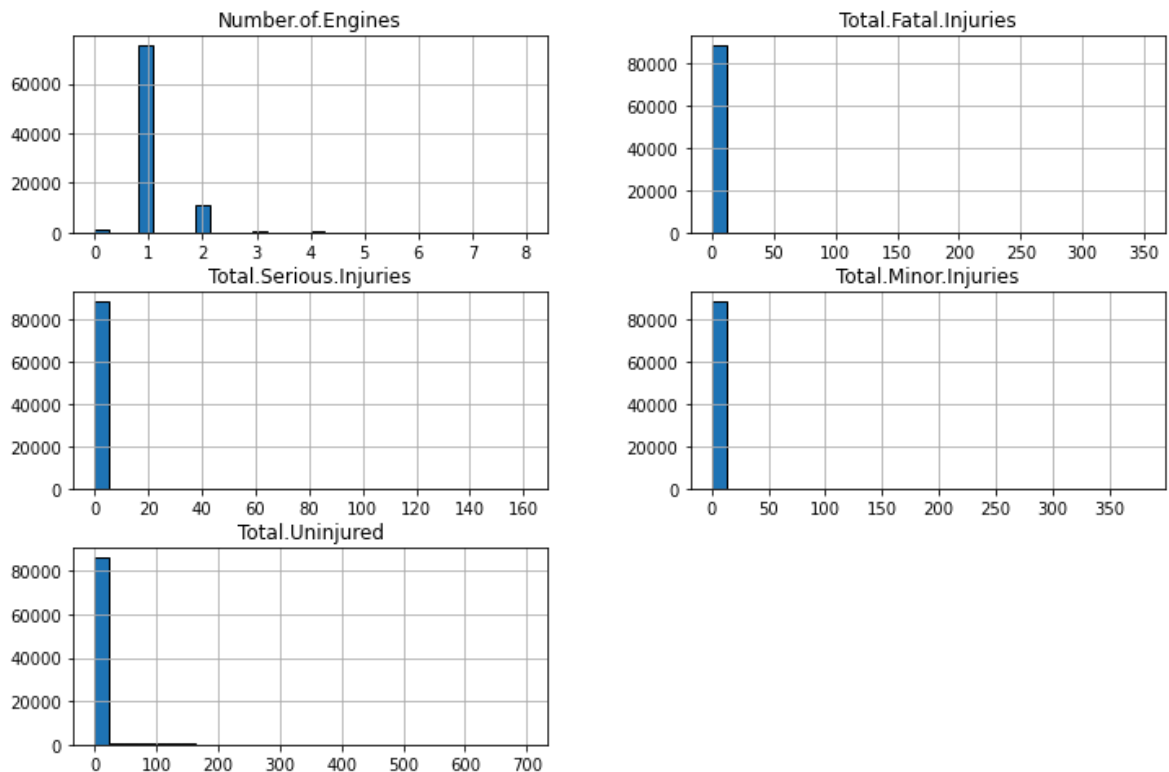


c) Histograms for Distribution

- They show how values are spread and highlight long tails caused by outliers.

```
In [21]: df[numeric_cols].hist(figsize=(12, 8), bins=30, edgecolor="black")
plt.suptitle("Distribution of Numeric Variables", fontsize=16)
plt.show()
```

Distribution of Numeric Variables



NB: I didn't remove outliers on my decision because they might be real accidents so if i drop them they might affect my analysis leading to biasness.

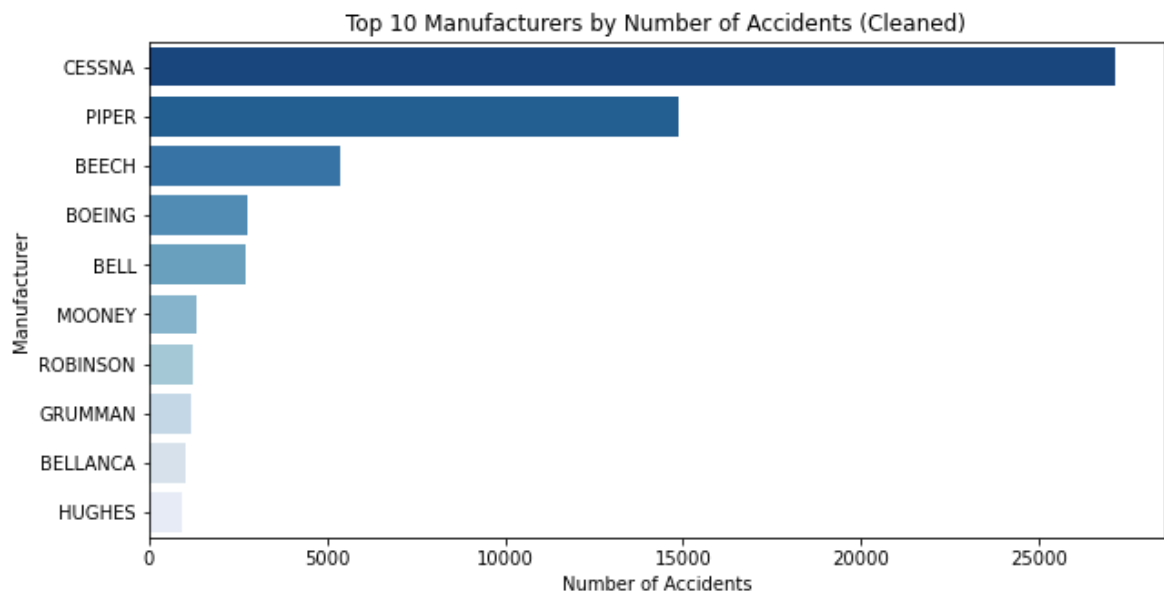
6: Data Visualization per Objective

Objective 1: Distribution & Severity by Model/Manufacturer

```
In [22]: # Clean the manufacturer names
df['Make_Cleaned'] = df['Make'].str.upper().str.strip()

# Count accidents by manufacturer
manufacturer_counts = df['Make_Cleaned'].value_counts().head(10)

# Plot
plt.figure(figsize=(10,5))
sns.barplot(x=manufacturer_counts.values, y=manufacturer_counts.index, palette="
plt.title("Top 10 Manufacturers by Number of Accidents (Cleaned)")
plt.xlabel("Number of Accidents")
plt.ylabel("Manufacturer")
plt.show()
```



Interpretation: The analysis of accident frequency by manufacturer indicates that a small number of aircraft producers account for the majority of reported accidents, with manufacturers like Cessna and Piper appearing most frequently in the dataset.

Adding severity columns that will help me to achieve my objective

```
In [23]: def classify_severity(row):
          if row['Total.Fatal.Injuries'] > 0:
              return "Fatal"
          elif row['Total.Serious.Injuries'] > 0:
              return "Serious"
          elif row['Total.Minor.Injuries'] > 0:
              return "Minor"
          else:
              return "No Injury"

          df["Injury_Severity"] = df.apply(classify_severity, axis=1)
```

Then group by it:

```
In [24]: severity_by_model = (
          df.groupby(['Make', 'Model'])['Injury_Severity']
            .value_counts()
            .unstack()
            .fillna(0)
          )
          severity_by_model.head()
```

Out[24]:

	Injury_Severity	Fatal	Minor	No Injury	Serious
Make	Model				
107.5 Flying Corporation	One Design DR 107	1.0	0.0	0.0	0.0
1200	G103	0.0	0.0	0.0	1.0
177MF LLC	PITTS MODEL 12	0.0	0.0	0.0	1.0
1977 Colfer-chan	STEEN SKYBOLT	0.0	1.0	0.0	0.0
1st Ftr Gp	FOCKE-WULF 190	1.0	0.0	0.0	0.0


Objective 2: Accident Characteristics & Fatal Incidents

In [25]:

```
# Create a "fatal" indicator
df['fatal_incident'] = np.where(df['Total.Fatal.Injuries'] > 0, 1, 0)

# Compare fatal vs non-fatal by damage level
sns.countplot(data=df, x='Aircraft.damage', hue='fatal_incident')
plt.title("Fatal vs Non-Fatal Incidents by Damage Level")
plt.show()

# Risk ratio by manufacturer
risk = df.groupby('Make')['fatal_incident'].mean().sort_values(ascending=False).
print(risk)
```



```
Make
107.5 Flying Corporation    1.0
Colliander                  1.0
Parrigin                    1.0
Coen                        1.0
Parkman                     1.0
Parker Warren               1.0
Parachute Icarus            1.0
Papa 51 Ltd., Co.           1.0
Pank                         1.0
Panaplane                   1.0
Name: fatal_incident, dtype: float64
```

Intepretation: Severe damage (e.g., "Destroyed") is strongly associated with fatal accidents while Minor/substantial damage tends to have mostly non-fatal outcomes. This indicates that aircraft damage level is a strong predictor of survivability.

Objective 3: Trends Over Time & Geography


```
In [26]: # Convert date
df['Event.Date'] = pd.to_datetime(df['Event.Date'], errors="coerce")

# Accidents per year
df['year'] = df['Event.Date'].dt.year
accidents_per_year = df.groupby('year').size()

plt.figure(figsize=(12,5))
sns.lineplot(x=accidents_per_year.index, y=accidents_per_year.values)
plt.title("Accident Frequency Over Time")
plt.xlabel("Year")
plt.ylabel("Number of Accidents")
plt.show()

# Geographical trends (using Plotly for interactivity)
accidents_by_country = df["Country"].value_counts().reset_index()
accidents_by_country.columns = ["Country", "Accident_Count"]

fig = px.choropleth(
    accidents_by_country,
    locations="Country",
    locationmode="country names",
    color="Accident_Count",
    title="Aviation Accident Frequency by Country"
)
fig.show()
```



Interpretation : The line chart provides insights into whether aviation safety has improved or worsened over time, Most often, you'll find a decline in accidents in recent decades

due to better aircraft engineering, pilot training, and regulations.

Also by countries, Darker-colored countries = higher number of accidents while Lighter-colored countries = lower accident counts. Usually, the United States dominates accident counts in such datasets (especially NTSB data) because of the large number of aircraft and reporting coverage but Some countries may appear low not because they're safer, but because they have fewer flights or less complete reporting.

7: Save the Cleaned Data and export it to tableau

```
In [27]: # Standardize manufacturer names  
df['Make'] = df['Make'].str.upper().str.strip()
```

```
In [28]: df.to_csv("Cleaned_AviationData.csv", index=False)
```

8: Conclusion

The study shows that accident distribution is concentrated among major manufacturers such as Cessna and Piper, largely reflecting their dominance in the market, but the likelihood of fatal outcomes differs significantly across aircraft types and models. Accident severity is closely tied to the extent of aircraft damage, with fatal incidents more common in cases of substantial or destroyed aircraft. Over time, accident frequencies demonstrate a general downward trend, reflecting improvements in aviation safety, though some spikes persist due to isolated events or reporting differences. Geographically, the United States records the highest number of accidents, consistent with its high flight volumes and strong reporting, while lower counts in other regions may reflect underreporting rather than inherently safer skies.

9: Recommendations

Based on the findings, the company should prioritize aircraft with lower fatality ratios rather than focusing solely on accident counts, and place emphasis on manufacturer-specific safety performance when making purchase decisions. Strengthening pilot training, routine maintenance, and safety audits is critical to reducing risk exposure, particularly for models with higher fatal accident proportions. Additionally, management should adopt a data-driven risk framework that integrates historical safety

records, geographical trends, and operational factors before expanding into aviation. Finally, ongoing monitoring and alignment with global safety standards will ensure that investment choices are both safe and sustainable in the long term.